

V3 VENTURES

Agentic OS — Requirement Brief

Central operating layer for the investment team · Version 3, 3 September 2026 · Prepared by Gargi, V3 Ventures

1. OBJECTIVE AND CONTEXT

V3 Ventures is a venture capital fund: we invest money in early-stage startups (our “portfolio companies”) and then track how each one performs, month after month, for years. Each company emails us a monthly pack of numbers. Today someone opens it, copies figures into a spreadsheet, and writes up what changed — for every company, every month.

We want to build an Agentic OS that becomes the central operating layer for V3’s investment team: data arrives from source systems on its own, the recurring analysis is drafted automatically, and the team’s time goes to judgement instead of assembly.

Scale to design for: 15–40 portfolio companies, each sending roughly one pack a month, plus board decks a few times a year. The system is used by V3’s investment team — a small number of internal users, not a public product.

2. DELIVERABLES AND CADENCE

What the system produces, and how often. Terms are explained in brackets at first use.

Deliverable	Cadence	What it contains	Primary source
Portfolio performance sheet	Monthly, automated, all companies	Each company’s numbers for the month, plus two separate layers of written commentary: objective commentary read out of the MIS (MIS = the monthly financial pack a company sends its investors — revenue, costs, cash) the way an analyst would read it, and subjective commentary drawn from recorded founder calls (what the founder said about traction, risks and priorities).	MIS files (OneDrive); founder call transcripts (Granola)
NAV report	Quarterly	NAV (NAV = net asset value, what V3’s stake in each company is currently worth) per company, added up to fund level (fund = the pool of money V3 invests from), plus what changed since last quarter and why.	MIS financials (OneDrive); ownership data (Affinity)
Performance dashboard	Real time — no manual updates	Live view across all companies: latest reported numbers, trend against earlier months, open flags, runway (runway = months of cash left before a company must raise again) and V3’s ownership stake. Refreshes as new files land rather than on a reporting cycle.	MIS financials (OneDrive); Affinity
Portfolio reports	On demand	Full portfolio report and single-company one-pagers on V3 templates, generated from live data and exportable to PDF, PPTX and XLSX.	MIS financials (OneDrive); Affinity; Granola

On the two commentary layers: objective commentary is what the numbers support — performance against the previous month and against plan, and what drove it. Subjective commentary is what the founder said on a call — priorities, concerns and asks. Both matter, and they must stay visibly separate so the reader always knows whether a statement is a measured number or a claim someone made.

3. DATA SOURCES AND INGESTION

Source	What it provides	Access method
OneDrive	Monthly MIS files and management accounts sent in by portfolio companies (the startups V3 has invested in).	Direct API pull. No manual uploads — this is the preferred design.
Affinity	V3’s CRM — company records, investment history, ownership, contacts.	API
Granola	Recordings and transcripts of calls with founders; the basis for all subjective commentary.	API
Claude	Reasoning, extraction and writing layer sitting across all of the above.	Connected to all sources

- **Primary design** — the system reads directly from source systems via API. Files that already sit in OneDrive are never re-uploaded by hand.
- **Fallback** — a separate backend ingestion layer with a form-based upload, for companies or documents that do not land in OneDrive. Fallback only, not the default path.
- **Onboarding** — a new portfolio company should be connected and producing structured output within 15 minutes, without engineering involvement.

4. INCOMING DATA, AND HOW IT MUST BE STANDARDISED

This is the hardest part of the build and the part most likely to be underestimated. Please read it before estimating.

- **No standard format.** Every company builds its own MIS. One puts revenue on a tab called “P&L”, the next calls it “Financials” and splits it by product line, a third sends a PDF export. Expect no two companies to match — and expect a company’s own format to change when it hires a new finance lead.
- **Mixed units.** Indian packs use lakh (100,000) and crore (10,000,000) alongside plain numbers, sometimes within one sheet. Some companies report in USD. Units must be detected explicitly, never inferred from magnitude.
- **Fiscal year runs April to March.** “Q1” means April–June, not January–March. Some companies report on calendar months regardless.
- **Revenue definitions vary.** Gross vs net of returns, GST-inclusive vs exclusive, GMV (total transaction value) vs recognised revenue. The same word means different things across two companies, and occasionally across two months at one company.

- **Restatements are normal.** A company sends March's numbers in April and revises them in May. The system must hold both and know which is current.
- **Missing is not zero.** Any field the system cannot find must be recorded as “not reported”, never as 0. This distinction matters more than it sounds — a zero silently corrupts every average and roll-up above it.

Required outcome — standardisation. However a company shapes its MIS, the system must map it into one V3 standard format: the field schema in Section 5. Everything downstream — dashboard, NAV, commentary, reports, Q&A — reads from that standard format and never from the original file. The original is retained only as the source a figure links back to. Standardising the incoming mess into a single shape is the core of this build, not a side effect of it.

Currency. Every figure is held and displayed in both INR crore and EUR. Companies report in INR (and occasionally USD); the system converts and stores both, recording the FX rate used and its date alongside each converted figure so any number can be reproduced later.

Corrections and attribution. The system has user logins. Any team member can correct an extracted figure; the correction is logged against that user and against the source file, and survives the same file being re-parsed. Accuracy is judged against the ingested files themselves — every figure must be checkable in one click against the cell it came from.

5. FIELDS TO EXTRACT

The V3 standard format every incoming MIS is mapped into. All monetary fields stored in INR crore and EUR. V3 will finalise this before build begins.

Group	Fields	Notes
Financial	Revenue (gross and net), cost of goods sold, gross margin, operating expenses by category, EBITDA, profit or loss	Monthly, with prior-month and plan comparison where the company provides a plan
Cash	Closing cash balance, monthly burn, runway	Runway is rarely stated and must be derived — closing cash divided by average burn over the last three months
Operating	Headcount, customers or users, orders or transactions, average order value, retention or churn	Varies by sector; not every company reports every field
Unit economics	CAC (cost to win a customer), contribution margin, payback period	Often absent. Extract where present; do not compute unless every input is available
Ownership	V3's ownership stake, latest round, post-money valuation, investment cost, cap table changes	Mostly from Affinity rather than the MIS
Qualitative	Founder-stated priorities, risks and asks	From call transcripts only, never inferred from the MIS

6. WORKFLOWS IN DETAIL

For each workflow: how it is done today, the automated workflow we want, and the output the system should produce.

Workflow	Current process	Desired agentic workflow	Expected output
Portfolio dashboard	Data sits in separate Excel trackers, decks and drive folders. A current view is assembled by hand before every review.	A live record per company that refreshes as new data arrives. No manual update step.	One dashboard — stage, ownership, valuation, NAV, MOIC and IRR (how many times our money has grown, and the annualised return), cash, burn (cash spent per month), runway, last funding round. Filterable, with click-through to the source file.
Automatic ingestion of files	Files arrive by email and are read manually. Figures are re-keyed into trackers, or Excel sheets dropped into Claude one at a time.	Files pulled from OneDrive via API as they land and mapped into V3's standard field schema, whatever shape the company sent. Anything the system is unsure of is flagged for a human, never guessed.	A standardised, timestamped record per company in INR crore and EUR, every field linked back to source file and cell. Ingestion status visible per company.
Company tracking rolled up to fund	Metrics tracked in separate company sheets. Fund-level totals rebuilt from scratch each cycle.	A standard metric set held per company over time and aggregated automatically to fund level, with restated figures reconciled against earlier submissions.	Time series per company — revenue, gross margin, burn, runway, headcount — and an always-current fund-level roll-up.
NAV and performance updates	NAV rebuilt manually in Excel each quarter. Valuation history lives in old file versions.	Valuation held per company on a consistent method and rolled to fund level quarterly, every change versioned and attributable.	Quarterly NAV by company and fund, MOIC and IRR, a bridge showing what moved value since last quarter, and full change history.
Monthly commentary	Written by an analyst after reading each MIS, if it gets written at all. Call context lives in personal notes.	System reads the MIS and drafts objective commentary; separately reads call transcripts and drafts subjective commentary. The two are never blended.	A monthly sheet per company — numbers, objective commentary and subjective commentary side by side, each statement traceable to the MIS line or transcript passage behind it.
Portfolio reports on demand	Assembled by hand from multiple sheets and decks. Takes days and is stale on arrival.	Generated on request from live data against fixed V3 templates. The analyst edits rather than assembles.	Full portfolio report and single-company one-pagers, produced in minutes.

Workflow	Current process	Desired agentic workflow	Expected output
Plain-language Q&A	Answered by opening files and reading. Cross-portfolio questions often go unasked because they cost too much time.	Ask anything across MIS files, transcripts and CRM records in plain English, with follow-ups. Every figure must come from a retrieved document; the system says the data is not available rather than estimating.	A direct answer citing source document and page — e.g. “which companies are spending more than ₹2 Cr a year on marketing?”
Risk and anomaly detection	Issues surface when someone happens to read a deck closely, or at the next board meeting.	Continuous scan of incoming files and transcripts against thresholds and prior periods, using flag categories V3 defines.	A ranked flag list with evidence — runway below threshold, missed targets, unusual movement in a line item, late or missing reporting, concerns raised on a call.
Cross-company comparison	Built ad hoc in Excel whenever a question comes up.	Any set of companies compared on a chosen metric set, normalised for stage and reporting period.	Comparison view and export — growth, gross margin, burn multiple, CAC and payback (cost to win a customer, and how long to earn it back), revenue per head — with peer benchmarking inside a sector.

7. KEY SUCCESS METRICS

Metric	Target
Portfolio data updated automatically	Target 90%+ of incoming reporting ingested and structured with no manual entry, and no manual uploads for companies reporting through OneDrive.
Accuracy of extracted numbers	Measured against the ingested files themselves. Near-zero tolerance for error on headline financials (revenue, cash, burn); every figure checkable in one click against the source cell, and correctable by a logged-in user.
Reduction in manual data-entry time	Target 80%+ less analyst time spent assembling data rather than interpreting it.
Time to generate a portfolio report	Minutes, not days — from request to a draft an analyst can edit.
Answers supported by correct citations	All of them. Every figure cited; where the data does not exist the system says so rather than estimating.
Time to onboard a new company	15 minutes end to end, from connecting the company's folder to first structured output.
Adoption and frequency of use	Everyday use by every member of the investment team, not a monthly reporting tool.
Risks and insights caught proactively	Substantially all flags caught against the categories V3 defines, with few enough false positives that the team keeps reading them.

Flag categories: to be defined by V3. Proposed starting set — runway below an agreed threshold, revenue below plan beyond an agreed band, gross margin falling, spend rising without matching revenue growth, late or missing monthly reporting, customer concentration shift, ownership or governance change, key person departure, and concerns raised by a founder on a call but absent from the MIS.